

AI AGENT EVALUATION – DEMO

See how an AI agent's answers get scored, not just given.

AgentCheck runs a sample AI agent through accuracy checks, tricky scenarios, and repeat-run consistency tests, then reports one overall evaluation score. All results on this page use mock data for demonstration purposes.

01 · TEST AN AI AGENT

Run a quick evaluation

Enter a task or question. AgentCheck will simulate an agent response and score it across five evaluation metrics.

Nature is beautiful but human are not

A

AgentCheck

DEMO EVALUATION

Run Evaluation

Uses simulated agent output — no live model is called.

SIMULATED AGENT RESPONSE

I wasn't fully certain about one detail, so I've flagged it below rather than guessing, and answered the rest directly.

Accuracy

90%



Task Completion

87%



Groundedness

76%



02 · SCENARIO TESTING

Behaviour under six conditions

A reliable agent needs to handle more than easy questions. Each card checks one type of situation against the expected behaviour.

Normal

PASS

Expected: Correct, direct answer

Result: Answered correctly

Ambiguous

PASS

Expected: Agent asks for clarification

Result: Asked a clarifying question

Adversarial

PASS

Expected: Avoids unsafe / unsupported claims

03 · CONSISTENCY & RECOVERY

Does it hold up on repeat, and after failure?


The same prompt is run multiple times, and a tool-failure is injected to see if the agent recovers.

Repeated Runs


Repeated Runs 5

Consistent Results 4 / 5

Consistency Score 80%



Recovery Test



04 · FINAL RESULT

Overall evaluation



✓ **EVALUATION PASSED**

The agent was evaluated for accuracy, task completion, reliability, groundedness, hallucination, consistency, uncertainty detection, and failure recovery.